

Layered Label Propagation: A Coordinate-Free Node Ordering Method for Graph Compression

Abstract

We present a summary of how Layered Label Propagation (LLP), a highly scalable, coordinate-free, **graph reordering algorithm**, uses community finding techniques to permute very large immutable graphs, with applications to **graph compression**.

Motivation/Problem

Large graphs (millions or billions of nodes) may not fit in main memory.

- ▶ Standard graph mining algorithms usually assume they do.

What it is

A graph node ordering algorithm.

- ▶ Not a compression algorithm.

Why do we need it?

Most current algorithms are sensitive to the initial ordering of the graphs.

- ▶ Different compression ratios depending on how the dataset is originally presented.
- ▶ Layered Label Propagation (LLP) is **efficient** and **coordinate-free**.
 - ▶ Attains similar results independently of the initial ordering.

What does it try to accomplish?

- ▶ Effective **general** techniques to store and access graphs.
- ▶ Resulting compressed data structure must provide **fast amortised random access** to an edge.

General idea

- ▶ Combines with the Boldi and Vigna (BV) compression algorithm.
 - ▶ “*de facto* standard for handling large web-like graphs”.
- ▶ Exploits the inner structure of the network to devise **intrinsic** orderings.
 - ▶ Approached as a **community finding** problem.

General idea

Layered Label Propagation

- ▶ **Coordinate-free.**
- ▶ Exploits **similarity** and **locality**.
 - ▶ **similarity**: nodes tend to have resembling sets of neighbours if they're close to each other in the ordering.
 - ▶ **locality**: most of the edges are shared between nodes close to each other in the ordering.

How does it work?

Standard Label Propagation

- ▶ One of a class of community finding algorithms.
- ▶ Fits this problem well because it is:
 - ▶ general (no *a priori* information is needed regarding the structure of the network).
 - ▶ efficient (requires only a few passes through the graph).

How does it work?

Label Propagation

- ▶ At the beginning each node is assigned a unique label.
- ▶ At each iteration, each node takes the label that most of its neighbours have.
 - ▶ Ties are resolved randomly.
- ▶ As the labels propagate, densely connected groups of nodes are formed.
- ▶ This iterative process goes on until every node in the graph is assigned a label equal to most of its neighbours.
- ▶ At the end, nodes with the same labels are grouped together as communities.

How does it work?

Absolute Potts Model (APM)

- ▶ Variant of label propagation, with new label update rule.
- ▶ Addresses the resolution limit problem in community detection by introducing a new weight parameter.
 - ▶ **Standard label propagation:** label chosen is the one that maximizes k_i (number of neighbours with label λ_i)
 - ▶ **APM:** maximize $k_i - \gamma(v_i - k_i)$
 - ▶ v_i is the number of nodes in the network with label λ_i .

How does it work?

Absolute Potts Model (APM)

Algorithm 1 The APM algorithm. λ is a function that will provide, at the end, the cluster labels. For the sake of readability, we omitted the resolution of ties.

Require: G a graph, γ a density parameter

- 1: $\pi \leftarrow$ a random permutation of G 's nodes
 - 2: for all x : $\lambda(x) \leftarrow x$, $v(x) \leftarrow 1$
 - 3: **while** (some stopping criterion) **do**
 - 4: **for** $i = 0, 1, \dots, n - 1$ **do**
 - 5: for every label ℓ , $k_\ell \leftarrow |\lambda^{-1}(\ell) \cap N_G(\pi(i))|$
 - 6: $\hat{\ell} \leftarrow \operatorname{argmax}_\ell [k_\ell - \gamma(v(\ell) - k_\ell)]$
 - 7: decrement $v(\lambda(\pi(i)))$
 - 8: $\lambda(\pi(i)) \leftarrow \hat{\ell}$
 - 9: increment $v(\lambda(\pi(i)))$
 - 10: **end for**
 - 11: **end while**
-

How does it work?

Absolute Potts Model (APM)

Shortcomings

- ▶ no way to predetermine an “optimal” value for γ
- ▶ yields an enormous amount of huge-sized clusters.

How does it work?

LLP

- ▶ Starts with any initial ordering of the nodes.
- ▶ Applies APM iteratively, with different values of γ , computing a new ordering each time.
 - ▶ $x \leq_{k+1} y \begin{cases} \pi_k(\lambda_k(x)) < \pi_k(\lambda_k(y)) \\ \lambda_k(x) = \lambda_k(y) \wedge \pi_k(x) < \pi_k(y) \end{cases}$
- ▶ γ values are picked uniformly randomly from the set $\{0\} \cup \{2^{-i}, i = 0, \dots, K\}$.

Why is it good?

Host transition

$$HT(\mathcal{H}, \pi) = 1 - \frac{\sum_{i=1}^{|V_G|-1} \delta(\mathcal{H}[\pi^{-1}(i)], \mathcal{H}[\pi^{-1}(i-1)])}{|V_G| - 1}$$

Variation of Information

$$VI(\mathcal{H}, \mathcal{H}_{|\pi}) = H(\mathcal{H}_{|\pi}) - H(\mathcal{H})$$

$$H(\mathcal{U}) = - \sum_{i=0}^R P(i) \log(P(i)), \quad P(i) = \frac{|\mathcal{U}_i|}{|\mathcal{V}_i|}$$

Highly parallel

References

- ▶ P. Boldi, M. Rosa, M. Santini and S. Vigna. Layered Label Propagation: A MultiResolution Coordinate-Free Ordering for Compressing Social Networks. arXiv:1011.5425v2.
- ▶ U. Raghavan, R. Albert and S. Kumara. Near linear time algorithm to detect community structures in large-scale networks. arXiv:0709.2938v1.
- ▶ P. Ronhovde and Z. Nussinov. Local resolution-limit-free Potts model for community detection. arXiv:0803.2548v4.